

CAPSTONE PROJECT 1:- Real Estate Case Study

Submitted by:- Ayub S Bijapur

```
In [1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
```

1.Import data

```
In [2]: df_train = pd.read_csv("train.csv")
df_train.head()
```

Out[2]:

	UID	BLOCKID	SUMLEVEL	COUNTYID	STATEID	state	state_ab	city	place
0	267822	NaN	140	53	36	New York	NY	Hamilton	Hamilton
1	246444	NaN	140	141	18	Indiana	IN	South Bend	Roseland
2	245683	NaN	140	63	18	Indiana	IN	Danville	Danville
3	279653	NaN	140	127	72	Puerto Rico	PR	San Juan	Guaynabo
4	247218	NaN	140	161	20	Kansas	KS	Manhattan	Manhattan City

5 rows × 80 columns

```
In [3]: df_test = pd.read_csv("test.csv")
df_test.head()
```

Out[3]:

	UID	BLOCKID	SUMLEVEL	COUNTYID	STATEID	state	state_ab	city	pla
0	255504	NaN	140	163	26	Michigan	MI	Detroit	Dearbc Heigl C
1	252676	NaN	140	1	23	Maine	ME	Auburn	Aubu C
2	276314	NaN	140	15	42	Pennsylvania	PA	Pine City	Millert
3	248614	NaN	140	231	21	Kentucky	KY	Monticello	Montice C
4	286865	NaN	140	355	48	Texas	TX	Corpus Christi	Edr

5 rows × 80 columns

```
In [4]: df_train.columns
```

```
Out[4]: Index(['UID', 'BLOCKID', 'SUMLEVEL', 'COUNTYID', 'STATEID', 'state',  
             'state_ab', 'city', 'place', 'type', 'primary', 'zip_code', 'area_code',  
             'lat', 'lng', 'ALand', 'AWater', 'pop', 'male_pop', 'female_pop',  
             'rent_mean', 'rent_median', 'rent_stdev', 'rent_sample_weight',  
             'rent_samples', 'rent_gt_10', 'rent_gt_15', 'rent_gt_20', 'rent_gt_25',  
             'rent_gt_30', 'rent_gt_35', 'rent_gt_40', 'rent_gt_50',  
             'universe_samples', 'used_samples', 'hi_mean', 'hi_median', 'hi_stdev',  
             'hi_sample_weight', 'hi_samples', 'family_mean', 'family_median',  
             'family_stdev', 'family_sample_weight', 'family_samples',  
             'hc_mortgage_mean', 'hc_mortgage_median', 'hc_mortgage_stdev',  
             'hc_mortgage_sample_weight', 'hc_mortgage_samples', 'hc_mean',  
             'hc_median', 'hc_stdev', 'hc_samples', 'hc_sample_weight',  
             'home_equity_second_mortgage', 'second_mortgage', 'home_equity', 'debt',  
             'second_mortgage_cdf', 'home_equity_cdf', 'debt_cdf', 'hs_degree',  
             'hs_degree_male', 'hs_degree_female', 'male_age_mean',  
             'male_age_median', 'male_age_stdev', 'male_age_sample_weight',  
             'male_age_samples', 'female_age_mean', 'female_age_median',  
             'female_age_stdev', 'female_age_sample_weight', 'female_age_samples',  
             'pct_own', 'married', 'married_snp', 'separated', 'divorced'],  
            dtype='object')
```

```
In [5]: df_test.columns
```

```
Out[5]: Index(['UID', 'BLOCKID', 'SUMLEVEL', 'COUNTYID', 'STATEID', 'state',  
             'state_ab', 'city', 'place', 'type', 'primary', 'zip_code', 'area_code',  
             'lat', 'lng', 'ALand', 'AWater', 'pop', 'male_pop', 'female_pop',  
             'rent_mean', 'rent_median', 'rent_stdev', 'rent_sample_weight',  
             'rent_samples', 'rent_gt_10', 'rent_gt_15', 'rent_gt_20', 'rent_gt_25',  
             'rent_gt_30', 'rent_gt_35', 'rent_gt_40', 'rent_gt_50',  
             'universe_samples', 'used_samples', 'hi_mean', 'hi_median', 'hi_stdev',  
             'hi_sample_weight', 'hi_samples', 'family_mean', 'family_median',  
             'family_stdev', 'family_sample_weight', 'family_samples',  
             'hc_mortgage_mean', 'hc_mortgage_median', 'hc_mortgage_stdev',  
             'hc_mortgage_sample_weight', 'hc_mortgage_samples', 'hc_mean',  
             'hc_median', 'hc_stdev', 'hc_samples', 'hc_sample_weight',  
             'home_equity_second_mortgage', 'second_mortgage', 'home_equity', 'debt',  
             'second_mortgage_cdf', 'home_equity_cdf', 'debt_cdf', 'hs_degree',  
             'hs_degree_male', 'hs_degree_female', 'male_age_mean',  
             'male_age_median', 'male_age_stdev', 'male_age_sample_weight',  
             'male_age_samples', 'female_age_mean', 'female_age_median',  
             'female_age_stdev', 'female_age_sample_weight', 'female_age_samples',  
             'pct_own', 'married', 'married_snp', 'separated', 'divorced'],  
            dtype='object')
```

```
In [6]: len(df_train)
```

```
Out[6]: 27321
```

```
In [7]: len(df_test)
```

```
Out[7]: 11709
```

```
In [8]: df_train.describe()
```

Out[8]:

	UID	BLOCKID	SUMLEVEL	COUNTYID	STATEID	zip_code	area_co
--	-----	---------	----------	----------	---------	----------	---------

count	27321.000000	0.0	27321.0	27321.000000	27321.000000	27321.000000	27321.0000
mean	257331.996303	NaN	140.0	85.646426	28.271806	50081.999524	596.5076
std	21343.859725	NaN	0.0	98.333097	16.392846	29558.115660	232.4974
min	220342.000000	NaN	140.0	1.000000	1.000000	602.000000	201.0000
25%	238816.000000	NaN	140.0	29.000000	13.000000	26554.000000	405.0000
50%	257220.000000	NaN	140.0	63.000000	28.000000	47715.000000	614.0000
75%	275818.000000	NaN	140.0	109.000000	42.000000	77093.000000	801.0000
max	294334.000000	NaN	140.0	840.000000	72.000000	99925.000000	989.0000

8 rows × 74 columns

In [9]: df_test.describe()

Out[9]:

	UID	BLOCKID	SUMLEVEL	COUNTYID	STATEID	zip_code	area_co
--	-----	---------	----------	----------	---------	----------	---------

count	11709.000000	0.0	11709.0	11709.000000	11709.000000	11709.000000	11709.0000
mean	257525.004783	NaN	140.0	85.710650	28.489196	50123.418396	593.5985
std	21466.372658	NaN	0.0	99.304334	16.607262	29775.134038	232.0742
min	220336.000000	NaN	140.0	1.000000	1.000000	601.000000	201.0000
25%	238819.000000	NaN	140.0	29.000000	13.000000	25570.000000	404.0000
50%	257651.000000	NaN	140.0	61.000000	28.000000	47362.000000	612.0000
75%	276300.000000	NaN	140.0	109.000000	42.000000	77406.000000	787.0000
max	294333.000000	NaN	140.0	810.000000	72.000000	99929.000000	989.0000

8 rows × 74 columns

In [10]: df_train.info()

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 27321 entries, 0 to 27320
```

```
Data columns (total 80 columns):
```

#	Column	Non-Null Count	Dtype
0	UID	27321 non-null	int64
1	BLOCKID	0 non-null	float64
2	SUMLEVEL	27321 non-null	int64
3	COUNTYID	27321 non-null	int64
4	STATEID	27321 non-null	int64
5	state	27321 non-null	object
6	state_ab	27321 non-null	object
7	city	27321 non-null	object
8	place	27321 non-null	object
9	type	27321 non-null	object
10	primary	27321 non-null	object
11	zip_code	27321 non-null	int64
12	area_code	27321 non-null	int64
13	lat	27321 non-null	float64
14	lng	27321 non-null	float64
15	ALand	27321 non-null	float64
16	AWater	27321 non-null	int64
17	pop	27321 non-null	int64
18	male_pop	27321 non-null	int64
19	female_pop	27321 non-null	int64
20	rent_mean	27007 non-null	float64
21	rent_median	27007 non-null	float64
22	rent_stdev	27007 non-null	float64
23	rent_sample_weight	27007 non-null	float64
24	rent_samples	27007 non-null	float64
25	rent_gt_10	27007 non-null	float64
26	rent_gt_15	27007 non-null	float64
27	rent_gt_20	27007 non-null	float64
28	rent_gt_25	27007 non-null	float64
29	rent_gt_30	27007 non-null	float64
30	rent_gt_35	27007 non-null	float64
31	rent_gt_40	27007 non-null	float64
32	rent_gt_50	27007 non-null	float64
33	universe_samples	27321 non-null	int64
34	used_samples	27321 non-null	int64
35	hi_mean	27053 non-null	float64
36	hi_median	27053 non-null	float64
37	hi_stdev	27053 non-null	float64
38	hi_sample_weight	27053 non-null	float64
39	hi_samples	27053 non-null	float64
40	family_mean	27023 non-null	float64
41	family_median	27023 non-null	float64
42	family_stdev	27023 non-null	float64
43	family_sample_weight	27023 non-null	float64
44	family_samples	27023 non-null	float64
45	hc_mortgage_mean	26748 non-null	float64
46	hc_mortgage_median	26748 non-null	float64
47	hc_mortgage_stdev	26748 non-null	float64
48	hc_mortgage_sample_weight	26748 non-null	float64
49	hc_mortgage_samples	26748 non-null	float64
50	hc_mean	26721 non-null	float64
51	hc_median	26721 non-null	float64
52	hc_stdev	26721 non-null	float64
53	hc_samples	26721 non-null	float64
54	hc_sample_weight	26721 non-null	float64
55	home_equity_second_mortgage	26864 non-null	float64
56	second_mortgage	26864 non-null	float64
57	home_equity	26864 non-null	float64
58	debt	26864 non-null	float64

```
59 second_mortgage_cdf      26864 non-null float64
60 home_equity_cdf          26864 non-null float64
61 debt_cdf                  26864 non-null float64
62 hs_degree                 27131 non-null float64
63 hs_degree_male            27121 non-null float64
64 hs_degree_female          27098 non-null float64
65 male_age_mean             27132 non-null float64
66 male_age_median           27132 non-null float64
67 male_age_stdev            27132 non-null float64
68 male_age_sample_weight    27132 non-null float64
69 male_age_samples          27132 non-null float64
70 female_age_mean           27115 non-null float64
71 female_age_median         27115 non-null float64
72 female_age_stdev          27115 non-null float64
73 female_age_sample_weight  27115 non-null float64
74 female_age_samples        27115 non-null float64
75 pct_own                   27053 non-null float64
76 married                   27130 non-null float64
77 married_snp               27130 non-null float64
78 separated                 27130 non-null float64
79 divorced                  27130 non-null float64
dtypes: float64(62), int64(12), object(6)
memory usage: 16.7+ MB
```

```
In [11]: df_test.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 11709 entries, 0 to 11708
```

```
Data columns (total 80 columns):
```

#	Column	Non-Null Count	Dtype
0	UID	11709 non-null	int64
1	BLOCKID	0 non-null	float64
2	SUMLEVEL	11709 non-null	int64
3	COUNTYID	11709 non-null	int64
4	STATEID	11709 non-null	int64
5	state	11709 non-null	object
6	state_ab	11709 non-null	object
7	city	11709 non-null	object
8	place	11709 non-null	object
9	type	11709 non-null	object
10	primary	11709 non-null	object
11	zip_code	11709 non-null	int64
12	area_code	11709 non-null	int64
13	lat	11709 non-null	float64
14	lng	11709 non-null	float64
15	ALand	11709 non-null	int64
16	AWater	11709 non-null	int64
17	pop	11709 non-null	int64
18	male_pop	11709 non-null	int64
19	female_pop	11709 non-null	int64
20	rent_mean	11561 non-null	float64
21	rent_median	11561 non-null	float64
22	rent_stdev	11561 non-null	float64
23	rent_sample_weight	11561 non-null	float64
24	rent_samples	11561 non-null	float64
25	rent_gt_10	11560 non-null	float64
26	rent_gt_15	11560 non-null	float64
27	rent_gt_20	11560 non-null	float64
28	rent_gt_25	11560 non-null	float64
29	rent_gt_30	11560 non-null	float64
30	rent_gt_35	11560 non-null	float64
31	rent_gt_40	11560 non-null	float64
32	rent_gt_50	11560 non-null	float64
33	universe_samples	11709 non-null	int64
34	used_samples	11709 non-null	int64
35	hi_mean	11587 non-null	float64
36	hi_median	11587 non-null	float64
37	hi_stdev	11587 non-null	float64
38	hi_sample_weight	11587 non-null	float64
39	hi_samples	11587 non-null	float64
40	family_mean	11573 non-null	float64
41	family_median	11573 non-null	float64
42	family_stdev	11573 non-null	float64
43	family_sample_weight	11573 non-null	float64
44	family_samples	11573 non-null	float64
45	hc_mortgage_mean	11441 non-null	float64
46	hc_mortgage_median	11441 non-null	float64
47	hc_mortgage_stdev	11441 non-null	float64
48	hc_mortgage_sample_weight	11441 non-null	float64
49	hc_mortgage_samples	11441 non-null	float64
50	hc_mean	11419 non-null	float64
51	hc_median	11419 non-null	float64
52	hc_stdev	11419 non-null	float64
53	hc_samples	11419 non-null	float64
54	hc_sample_weight	11419 non-null	float64
55	home_equity_second_mortgage	11489 non-null	float64
56	second_mortgage	11489 non-null	float64
57	home_equity	11489 non-null	float64
58	debt	11489 non-null	float64

```

59 second_mortgage_cdf      11489 non-null float64
60 home_equity_cdf         11489 non-null float64
61 debt_cdf                11489 non-null float64
62 hs_degree               11624 non-null float64
63 hs_degree_male          11620 non-null float64
64 hs_degree_female        11604 non-null float64
65 male_age_mean           11625 non-null float64
66 male_age_median         11625 non-null float64
67 male_age_stdev          11625 non-null float64
68 male_age_sample_weight  11625 non-null float64
69 male_age_samples        11625 non-null float64
70 female_age_mean         11613 non-null float64
71 female_age_median       11613 non-null float64
72 female_age_stdev        11613 non-null float64
73 female_age_sample_weight 11613 non-null float64
74 female_age_samples      11613 non-null float64
75 pct_own                 11587 non-null float64
76 married                 11625 non-null float64
77 married_snp             11625 non-null float64
78 separated               11625 non-null float64
79 divorced                11625 non-null float64
dtypes: float64(61), int64(13), object(6)
memory usage: 7.1+ MB

```

2. Figure out the primary key and look for the requirement of indexing

UID is unique userID value in the train and test dataset. So an index can be created from the UID feature

```

In [12]: #Set the DataFrame index using existing columns.
df_train.set_index(keys=['UID'],inplace=True)
df_test.set_index(keys=['UID'],inplace=True)

```

```

In [13]: df_train.head(2)

```

```

Out[13]:
```

	BLOCKID	SUMLEVEL	COUNTYID	STATEID	state	state_ab	city	place	type
UID									
267822	NaN	140	53	36	New York	NY	Hamilton	Hamilton	City
246444	NaN	140	141	18	Indiana	IN	South Bend	Roseland	City

2 rows × 79 columns

```

In [14]: df_test.head(2)

```

Out[14]:

	BLOCKID	SUMLEVEL	COUNTYID	STATEID	state	state_ab	city	place	type
UID									
255504	NaN	140	163	26	Michigan	MI	Detroit	Dearborn Heights City	CDP
252676	NaN	140	1	23	Maine	ME	Auburn	Auburn City	City

2 rows × 79 columns

3. Gauge the fill rate of the variables and devise plans for missing value treatment. Please explain explicitly the reason for the treatment chosen for each variable.

```
In [15]: #percentage of missing values in train set
missing_list_train=df_train.isnull().sum() *100/len(df_train)
missing_values_df_train=pd.DataFrame(missing_list_train,columns=['Percentage of missing values'])
missing_values_df_train.sort_values(by=['Percentage of missing values'],inplace=True)
missing_values_df_train[missing_values_df_train['Percentage of missing values'] >0]
#BLOCKID can be dropped, since it is 100%missing values
```

Out[15]:

Percentage of missing values	
BLOCKID	100.000000
hc_samples	2.196113
hc_mean	2.196113
hc_median	2.196113
hc_stdev	2.196113
hc_sample_weight	2.196113
hc_mortgage_mean	2.097288
hc_mortgage_stdev	2.097288
hc_mortgage_sample_weight	2.097288
hc_mortgage_samples	2.097288

```
In [16]: #percentage of missing values in test set
missing_list_test=df_test.isnull().sum() *100/len(df_train)
missing_values_df_test=pd.DataFrame(missing_list_test,columns=['Percentage of missing values'])
missing_values_df_test.sort_values(by=['Percentage of missing values'],inplace=True)
missing_values_df_test[missing_values_df_test['Percentage of missing values'] >0]
#BLOCKID can be dropped, since it is 43%missing values
```


Out[16]:

Percentage of missing values	
BLOCKID	42.857143
hc_samples	1.061455
hc_mean	1.061455
hc_median	1.061455
hc_stdev	1.061455
hc_sample_weight	1.061455
hc_mortgage_mean	0.980930
hc_mortgage_stdev	0.980930
hc_mortgage_sample_weight	0.980930
hc_mortgage_samples	0.980930

```
In [17]: df_train.drop(columns=['BLOCKID', 'SUMLEVEL'], axis=1, inplace=True)
#SUMLEVEL does not have any predictive power and no variance
```

```
In [18]: df_test .drop(columns=['BLOCKID', 'SUMLEVEL'], axis=1, inplace=True)
#SUMLEVEL does not have any predictive power
```

```
In [19]: # Imputing missing values with mean
missing_train_cols=[]
for col in df_train.columns:
    if df_train[col].isna().sum() !=0:
        missing_train_cols.append(col)
print(missing_train_cols)

['rent_mean', 'rent_median', 'rent_stdev', 'rent_sample_weight', 'rent_samples',
'rent_gt_10', 'rent_gt_15', 'rent_gt_20', 'rent_gt_25', 'rent_gt_30', 'rent_gt_35',
'rent_gt_40', 'rent_gt_50', 'hi_mean', 'hi_median', 'hi_stdev', 'hi_sample_weight',
'hi_samples', 'family_mean', 'family_median', 'family_stdev', 'family_sample_weight',
'family_samples', 'hc_mortgage_mean', 'hc_mortgage_median', 'hc_mortgage_stdev',
'hc_mortgage_sample_weight', 'hc_mortgage_samples', 'hc_mean', 'hc_median',
'hc_stdev', 'hc_samples', 'hc_sample_weight', 'home_equity_second_mortgage',
'second_mortgage', 'home_equity', 'debt', 'second_mortgage_cdf', 'home_equity_cdf',
'debt_cdf', 'hs_degree', 'hs_degree_male', 'hs_degree_female', 'male_age_mean',
'male_age_median', 'male_age_stdev', 'male_age_sample_weight', 'male_age_samples',
'female_age_mean', 'female_age_median', 'female_age_stdev', 'female_age_sample_weight',
'female_age_samples', 'pct_own', 'married', 'married_snp', 'separated',
'divorced']
```

```
In [20]: # Imputing missing values with mean
missing_test_cols=[]
for col in df_test.columns:
    if df_test[col].isna().sum() !=0:
        missing_test_cols.append(col)
print(missing_test_cols)
```

```
[ 'rent_mean', 'rent_median', 'rent_stdev', 'rent_sample_weight', 'rent_samples',
  'rent_gt_10', 'rent_gt_15', 'rent_gt_20', 'rent_gt_25', 'rent_gt_30', 'rent_gt_35',
  'rent_gt_40', 'rent_gt_50', 'hi_mean', 'hi_median', 'hi_stdev', 'hi_sample_weight',
  'hi_samples', 'family_mean', 'family_median', 'family_stdev', 'family_sample_weight',
  'family_samples', 'hc_mortgage_mean', 'hc_mortgage_median', 'hc_mortgage_stdev',
  'hc_mortgage_sample_weight', 'hc_mortgage_samples', 'hc_mean', 'hc_median',
  'hc_stdev', 'hc_samples', 'hc_sample_weight', 'home_equity_second_mortgage',
  'second_mortgage', 'home_equity', 'debt', 'second_mortgage_cdf', 'home_equity_cdf',
  'debt_cdf', 'hs_degree', 'hs_degree_male', 'hs_degree_female', 'male_age_mean',
  'male_age_median', 'male_age_stdev', 'male_age_sample_weight', 'male_age_samples',
  'female_age_mean', 'female_age_median', 'female_age_stdev', 'female_age_sample_weight',
  'female_age_samples', 'pct_own', 'married', 'married_snp', 'separated', 'divorced']
```

```
In [21]: # Missing cols are all numerical variables
for col in df_train.columns:
    if col in (missing_train_cols):
        df_train[col].replace(np.nan, df_train[col].mean(), inplace=True)
```

```
In [22]: # Missing cols are all numerical variables
for col in df_test.columns:
    if col in (missing_test_cols):
        df_test[col].replace(np.nan, df_test[col].mean(), inplace=True)
```

```
In [23]: df_train.isna().sum().sum()
```

```
Out[23]: 0
```

```
In [24]: df_test.isna().sum().sum()
```

```
Out[24]: 0
```

Exploratory Data Analysis (EDA):

4.Perform debt analysis. You may take the following steps:

a) Explore the top 2,500 locations where the percentage of households with a second mortgage is the highest and percent ownership is above 10 percent. Visualize using geo-map. You may keep the upper limit for the percent of households with a second mortgage to 50 percent

```
In [25]: from pandasql import sqldf
q1 = "select place,pct_own,second_mortgage,lat,lng from df_train where pct_own >0.1"
pysqldf = lambda q: sqldf(q, globals())
df_train_location_mort_pct=pysqldf(q1)
```

```
In [26]: df_train_location_mort_pct.head()
```

Out[26]:

	place	pct_own	second_mortgage	lat	lng
0	Worcester City	0.20247	0.43363	42.254262	-71.800347
1	Harbor Hills	0.15618	0.31818	40.751809	-73.853582
2	Glen Burnie	0.22380	0.30212	39.127273	-76.635265
3	Egypt Lake-leto	0.11618	0.28972	28.029063	-82.495395
4	Lincolnwood	0.14228	0.28899	41.967289	-87.652434


```
In [27]: import plotly.express as px
import plotly.graph_objects as go
```

```
In [28]: fig = go.Figure(data=go.Scattergeo(
    lat = df_train_location_mort_pct['lat'],
    lon = df_train_location_mort_pct['lng']),
)
fig.update_layout(
    geo=dict(
        scope = 'north america',
        showland = True,
        landcolor = "rgb(212, 212, 212)",
        subunitcolor = "rgb(255, 255, 255)",
        countrycolor = "rgb(255, 255, 255)",
        showlakes = True,
        lakecolor = "rgb(255, 255, 255)",
        showsubunits = True,
        showcountries = True,
        resolution = 50,
        projection = dict(
            type = 'conic conformal',
            rotation_lon = -100
        ),
        lonaxis = dict(
            showgrid = True,
            gridwidth = 0.5,
            range= [ -140.0, -55.0 ],
            dtick = 5
        ),
        lataxis = dict (
            showgrid = True,
            gridwidth = 0.5,
            range= [ 20.0, 60.0 ],
            dtick = 5
        )
    ),
    title='Top 2,500 locations with second mortgage is the highest and percent owned',
fig.show()
```

b) Use the following bad debt equation: $\text{Bad Debt} = \text{P (Second Mortgage} \cap \text{Home Equity Loan)}$ $\text{Bad Debt} = \text{second_mortgage} + \text{home_equity} - \text{home_equity_second_mortgage}$ c) Create pie charts to show over

```
In [29]: df_train['bad_debt']=df_train['second_mortgage']+df_train['home_equity']-df_train[

In [30]: df_train['bins'] = pd.cut(df_train['bad_debt'],bins=[0,0.10,1], labels=["less than
df_train.groupby(['bins']).size().plot(kind='pie',subplots=True,startangle=90, auto
plt.axis('equal')
plt.show()
```



d) Create Box and whisker plot and analyze the distribution for 2nd mortgage, home equity, good debt, and bad debt for different cities

```
In [31]: cols=[]
df_train.columns
```


```
Out[31]: Index(['COUNTYID', 'STATEID', 'state', 'state_ab', 'city', 'place', 'type',
'primary', 'zip_code', 'area_code', 'lat', 'lng', 'ALand', 'AWater',
'pop', 'male_pop', 'female_pop', 'rent_mean', 'rent_median',
'rent_stdev', 'rent_sample_weight', 'rent_samples', 'rent_gt_10',
'rent_gt_15', 'rent_gt_20', 'rent_gt_25', 'rent_gt_30', 'rent_gt_35',
'rent_gt_40', 'rent_gt_50', 'universe_samples', 'used_samples',
'hi_mean', 'hi_median', 'hi_stdev', 'hi_sample_weight', 'hi_samples',
'family_mean', 'family_median', 'family_stdev', 'family_sample_weight',
'family_samples', 'hc_mortgage_mean', 'hc_mortgage_median',
'hc_mortgage_stdev', 'hc_mortgage_sample_weight', 'hc_mortgage_samples',
'hc_mean', 'hc_median', 'hc_stdev', 'hc_samples', 'hc_sample_weight',
'home_equity_second_mortgage', 'second_mortgage', 'home_equity', 'debt',
'second_mortgage_cdf', 'home_equity_cdf', 'debt_cdf', 'hs_degree',
'hs_degree_male', 'hs_degree_female', 'male_age_mean',
'male_age_median', 'male_age_stdev', 'male_age_sample_weight',
'male_age_samples', 'female_age_mean', 'female_age_median',
'female_age_stdev', 'female_age_sample_weight', 'female_age_samples',
'pct_own', 'married', 'married_snp', 'separated', 'divorced',
'bad_debt', 'bins'],
dtype='object')
```

```
In [32]: #Taking Hamilton and Manhattan cities data
cols=['second_mortgage', 'home_equity', 'debt', 'bad_debt']
df_box_hamilton=df_train.loc[df_train['city'] == 'Hamilton']
df_box_manhattan=df_train.loc[df_train['city'] == 'Manhattan']
df_box_city=pd.concat([df_box_hamilton, df_box_manhattan])
df_box_city.head(4)
```

Out [32]:


	COUNTYID	STATEID	state	state_ab	city	place	type	primary	zip_code
UID									
267822	53	36	New York	NY	Hamilton	Hamilton	City	tract	13346
263797	21	34	New Jersey	NJ	Hamilton	Yardville	City	tract	8610
270979	17	39	Ohio	OH	Hamilton	Hamilton City	Village	tract	45015
259028	95	28	Mississippi	MS	Hamilton	Hamilton	CDP	tract	39746

4 rows × 79 columns




In [33]:

```
plt.figure(figsize=(10,5))
sns.boxplot(data=df_box_city,x='second_mortgage', y='city',width=0.5,palette="Set3")
plt.show()
```




In [34]:


```
plt.figure(figsize=(10,5))
sns.boxplot(data=df_box_city,x='home_equity', y='city',width=0.5,palette="Set3")
plt.show()
```



```
In [35]: plt.figure(figsize=(10,5))  
sns.boxplot(data=df_box_city,x='debt', y='city',width=0.5,palette="Set3")  
plt.show()
```



```
In [36]: plt.figure(figsize=(10,5))  
sns.boxplot(data=df_box_city,x='bad_debt', y='city',width=0.5,palette="Set3")  
plt.show()
```




Manhattan has higher metrics compared to Hamilton

e) Create a collated income distribution chart for family income, house hold income, and remaining income

```
In [37]: sns.distplot(df_train['hi_mean'])  
plt.title('Household income distribution chart')  
plt.show()
```

/home/spx072/anaconda3/lib/python3.10/site-packages/seaborn/distributions.py:2619:
FutureWarning:


`distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).



```
In [38]: sns.distplot(df_train['family_mean'])  
plt.title('Family income distribution chart')  
plt.show()
```

/home/spx072/anaconda3/lib/python3.10/site-packages/seaborn/distributions.py:2619:
FutureWarning:


`distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).



```
In [39]: sns.distplot(df_train['family_mean']-df_train['hi_mean'])  
plt.title('Remaining income distribution chart')  
plt.show()
```

/home/spx072/anaconda3/lib/python3.10/site-packages/seaborn/distributions.py:2619:
FutureWarning:

`distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).



Exploratory Data Analysis (EDA):

1. Perform EDA and come out with insights into population density and age. You may have to derive new fields (make sure to weight averages for accurate measurements):

```
In [40]: #plt.figure(figsize=(25,10))
fig,(ax1,ax2,ax3)=plt.subplots(3,1)
sns.distplot(df_train['pop'],ax=ax1)
sns.distplot(df_train['male_pop'],ax=ax2)
sns.distplot(df_train['female_pop'],ax=ax3)
plt.subplots_adjust(wspace=0.8,hspace=0.8)
plt.tight_layout()
plt.show()
```

/home/spx072/anaconda3/lib/python3.10/site-packages/seaborn/distributions.py:2619: FutureWarning:


`distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

/home/spx072/anaconda3/lib/python3.10/site-packages/seaborn/distributions.py:2619: FutureWarning:

`distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

/home/spx072/anaconda3/lib/python3.10/site-packages/seaborn/distributions.py:2619: FutureWarning:

`distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).




```
In [41]: #plt.figure(figsize=(25,10))
fig,(ax1,ax2)=plt.subplots(2,1)
sns.distplot(df_train['male_age_mean'],ax=ax1)
sns.distplot(df_train['female_age_mean'],ax=ax2)
plt.subplots_adjust(wspace=0.8,hspace=0.8)
plt.tight_layout()
plt.show()
```

```
/home/spx072/anaconda3/lib/python3.10/site-packages/seaborn/distributions.py:2619:
FutureWarning:
```

``distplot`` is a deprecated function and will be removed in a future version. Please adapt your code to use either ``displot`` (a figure-level function with similar flexibility) or ``histplot`` (an axes-level function for histograms).

```
/home/spx072/anaconda3/lib/python3.10/site-packages/seaborn/distributions.py:2619:
FutureWarning:
```

``distplot`` is a deprecated function and will be removed in a future version. Please adapt your code to use either ``displot`` (a figure-level function with similar flexibility) or ``histplot`` (an axes-level function for histograms).



a) Use pop and ALand variables to create a new field called population density


```
In [42]: df_train['pop_density']=df_train['pop']/df_train['ALand']
```

```
In [43]: df_test['pop_density']=df_test['pop']/df_test['ALand']
```

```
In [44]: sns.distplot(df_train['pop_density'])
plt.title('Population Density')
plt.show() # Very less density is noticed
```

```
/home/spx072/anaconda3/lib/python3.10/site-packages/seaborn/distributions.py:2619:
FutureWarning:
```

``distplot`` is a deprecated function and will be removed in a future version. Please adapt your code to use either ``displot`` (a figure-level function with similar flexibility) or ``histplot`` (an axes-level function for histograms).



b) Use `male_age_median`, `female_age_median`, `male_pop`, and `female_pop` to create a new field called `median age` c) Visualize the findings using appropriate chart type

```
In [45]: df_train['age_median']=(df_train['male_age_median']+df_train['female_age_median']),
df_test['age_median']=(df_test['male_age_median']+df_test['female_age_median'])/2
```

```
In [46]: df_train[['male_age_median', 'female_age_median', 'male_pop', 'female_pop', 'age_median',
```


```
Out[46]:      male_age_median  female_age_median  male_pop  female_pop  age_median
```

UID					
267822	44.00000	45.33333	2612	2618	44.666665
246444	32.00000	37.58333	1349	1284	34.791665
245683	40.83333	42.83333	3643	3238	41.833330
279653	48.91667	50.58333	1141	1559	49.750000
247218	22.41667	21.58333	2586	3051	22.000000

```
In [47]: sns.distplot(df_train['age_median'])
plt.title('Median Age')
plt.show()
# Age of population is mostly between 20 and 60
# Majority are of age around 40
# Median age distribution has a gaussian distribution
# Some right skewness is noticed
```

```
/home/spx072/anaconda3/lib/python3.10/site-packages/seaborn/distributions.py:2619: FutureWarning:
```


`distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).



```
In [48]: sns.boxplot(df_train['age_median'])  
plt.title('Population Density')  
plt.show()
```

```
/home/spx072/anaconda3/lib/python3.10/site-packages/seaborn/_decorators.py:36: FutureWarning:
```

Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.



2. Create bins for population into a new variable by selecting appropriate class interval so that the number of categories don't exceed 5 for the ease of analysis.

```
In [49]: df_train['pop'].describe()
```

```
Out[49]: count    27321.000000
mean      4316.032685
std       2169.226173
min        0.000000
25%       2885.000000
50%       4042.000000
75%       5430.000000
max      53812.000000
Name: pop, dtype: float64
```

```
In [50]: df_train['pop_bins']=pd.cut(df_train['pop'],bins=5,labels=['very low','low','medium','high','very high'])
```

```
In [51]: df_train[['pop','pop_bins']]
```


Out[51]:

	pop	pop_bins
UID		
267822	5230	very low
246444	2633	very low
245683	6881	very low
279653	2700	very low
247218	5637	very low
...
279212	1847	very low
277856	4155	very low
233000	2829	very low
287425	11542	low
265371	3726	very low

27321 rows × 2 columns

In [52]: `df_train['pop_bins'].value_counts()`

Out[52]:

very low	27058
low	246
medium	9
high	7
very high	1

Name: pop_bins, dtype: int64

a) Analyze the married, separated, and divorced population for these population brackets

In [53]: `df_train.groupby(by='pop_bins')[['married','separated','divorced']].count()`

Out[53]:

	married	separated	divorced
pop_bins			

pop_bins			
very low	27058	27058	27058
low	246	246	246
medium	9	9	9
high	7	7	7
very high	1	1	1

In [54]: `df_train.groupby(by='pop_bins')[['married','separated','divorced']].agg(["mean", "r`

Out[54]:

	married		separated		divorced	
	mean	median	mean	median	mean	median
pop_bins						
very low	0.507548	0.524680	0.019126	0.013650	0.100504	0.096020
low	0.584894	0.593135	0.015833	0.011195	0.075348	0.070045
medium	0.655737	0.618710	0.005003	0.004120	0.065927	0.064890
high	0.503359	0.335660	0.008141	0.002500	0.039030	0.010320
very high	0.734740	0.734740	0.004050	0.004050	0.030360	0.030360


1. Very high population group has more married people and less percentage of separated and divorced couples

2. In very low population groups, there are more divorced people

b) Visualize using appropriate chart type

```
In [55]: plt.figure(figsize=(10,5))
pop_bin_married=df_train.groupby(by='pop_bins')[['married','separated','divorced']]
pop_bin_married.plot(figsize=(20,8))
plt.legend(loc='best')
plt.show()
```

<Figure size 1000x500 with 0 Axes>



3. Please detail your observations for rent as a percentage of income at an overall level, and for different states.

```
In [56]: rent_state_mean=df_train.groupby(by='state')['rent_mean'].agg(["mean"])
rent_state_mean.head()
```

Out[56]:

	mean
state	
Alabama	774.004927
Alaska	1185.763570
Arizona	1097.753511
Arkansas	720.918575
California	1471.133857

In [57]: `income_state_mean=df_train.groupby(by='state')['family_mean'].agg(["mean"])`
`income_state_mean.head()`

Out[57]:

	mean
state	
Alabama	67030.064213
Alaska	92136.545109
Arizona	73328.238798
Arkansas	64765.377850
California	87655.470820

In [58]: `rent_perc_of_income=rent_state_mean['mean']/income_state_mean['mean']`
`rent_perc_of_income.head(10)`

Out[58]:

state	
Alabama	0.011547
Alaska	0.012870
Arizona	0.014970
Arkansas	0.011131
California	0.016783
Colorado	0.013529
Connecticut	0.012637
Delaware	0.012929
District of Columbia	0.013198
Florida	0.015772

Name: mean, dtype: float64

In [59]: `#overall level rent as a percentage of income`
`sum(df_train['rent_mean'])/sum(df_train['family_mean'])`

Out[59]: 0.013358170721473864


4. Perform correlation analysis for all the relevant variables by creating a heatmap. Describe your findings.

In [60]: `df_train.columns`

```
Out[60]: Index(['COUNTYID', 'STATEID', 'state', 'state_ab', 'city', 'place', 'type',
'primary', 'zip_code', 'area_code', 'lat', 'lng', 'ALand', 'AWater',
'pop', 'male_pop', 'female_pop', 'rent_mean', 'rent_median',
'rent_stddev', 'rent_sample_weight', 'rent_samples', 'rent_gt_10',
'rent_gt_15', 'rent_gt_20', 'rent_gt_25', 'rent_gt_30', 'rent_gt_35',
'rent_gt_40', 'rent_gt_50', 'universe_samples', 'used_samples',
'hi_mean', 'hi_median', 'hi_stddev', 'hi_sample_weight', 'hi_samples',
'family_mean', 'family_median', 'family_stddev', 'family_sample_weight',
'family_samples', 'hc_mortgage_mean', 'hc_mortgage_median',
'hc_mortgage_stddev', 'hc_mortgage_sample_weight', 'hc_mortgage_samples',
'hc_mean', 'hc_median', 'hc_stddev', 'hc_samples', 'hc_sample_weight',
'home_equity_second_mortgage', 'second_mortgage', 'home_equity', 'debt',
'second_mortgage_cdf', 'home_equity_cdf', 'debt_cdf', 'hs_degree',
'hs_degree_male', 'hs_degree_female', 'male_age_mean',
'male_age_median', 'male_age_stddev', 'male_age_sample_weight',
'male_age_samples', 'female_age_mean', 'female_age_median',
'female_age_stddev', 'female_age_sample_weight', 'female_age_samples',
'pct_own', 'married', 'married_snp', 'separated', 'divorced',
'bad_debt', 'bins', 'pop_density', 'age_median', 'pop_bins'],
dtype='object')
```

```
In [61]: cor=df_train[['COUNTYID','STATEID','zip_code','type','pop', 'family_mean',
'second_mortgage', 'home_equity', 'debt','hs_degree',
'age_median','pct_own', 'married','separated', 'divorced']].corr()
```

```
In [62]: plt.figure(figsize=(20,10))
sns.heatmap(cor,annot=True,cmap='coolwarm')
plt.show()
```



1.High positive correaltion is noticed between pop, male_pop and female_pop

2.High positive correaltion is noticed between rent_mean,hi_mean, family_mean,hc_mean

Data Pre-processing:

1. The economic multivariate data has a significant number of measured variables. The goal is to find where the measured variables depend on a number of smaller unobserved common factors or latent variables. 2. Each variable is assumed to be dependent upon a linear combination of the

common factors, and the coefficients are known as loadings. Each measured variable also includes a component due to independent random variability, known as "specific variance" because it is specific to one variable. Obtain the common factors and then plot the loadings. Use factor analysis to find latent variables in our dataset and gain insight into the linear relationships in the data. Following are the list of latent variables:

- Highschool graduation rates
- Median population age
- Second mortgage statistics
- Percent own
- Bad debt expense

```
In [63]: from sklearn.decomposition import FactorAnalysis  
from factor_analyzer import FactorAnalyzer
```

```
In [64]: fa=FactorAnalyzer(n_factors=5)  
fa.fit_transform(df_train.select_dtypes(exclude= ('object', 'category')))  
fa.loadings_
```

```

Out[64]: array([[ -1.12589168e-01,  1.95646473e-02, -2.39331081e-02,
                -6.27632630e-02,  4.23474744e-02],
                [ -1.10186764e-01,  1.33506220e-02,  2.79651243e-02,
                -1.49825864e-01,  1.10838806e-01],
                [ -8.28678626e-02,  5.16372371e-02, -1.36451864e-01,
                -4.98918600e-02, -1.04024838e-01],
                [  1.80961136e-02,  1.92013749e-02,  5.81329912e-03,
                2.64842745e-02, -6.12442620e-03],
                [  9.02324727e-02, -9.72544310e-02, -6.54601334e-02,
                -1.33145905e-01, -1.48594602e-01],
                [ -1.07335697e-02, -4.12376813e-02,  1.45853480e-01,
                8.80433214e-03,  1.08227564e-01],
                [ -4.28796975e-02, -2.09780214e-02,  3.66726857e-02,
                -9.45597367e-02,  5.91380506e-02],
                [ -2.44243043e-03, -1.53245408e-02, -2.68300838e-03,
                -4.52473026e-02,  2.37240645e-02],
                [  7.92164309e-02,  9.57453307e-01, -8.71151634e-02,
                -6.59923727e-03, -3.97273224e-02],
                [  7.39808182e-02,  9.18750508e-01, -1.08834838e-01,
                -2.79371563e-02, -3.93153685e-02],
                [  8.06598878e-02,  9.47839209e-01, -6.08006517e-02,
                1.53627100e-02, -3.86977301e-02],
                [  7.70052115e-01,  9.84675280e-03, -3.71249731e-02,
                1.14949040e-01, -1.23784691e-01],
                [  7.18615876e-01,  6.24980344e-03, -4.59787391e-02,
                1.09109690e-01, -1.35301916e-01],
                [  7.07647251e-01,  2.46625398e-02, -1.00860886e-02,
                1.04472489e-01,  7.72381273e-02],
                [ -1.34545479e-01,  3.36809304e-01, -4.87894982e-01,
                -4.15446298e-02,  3.17608549e-01],
                [  2.31079719e-01,  4.37729796e-01, -6.40209229e-01,
                -2.52311077e-02,  3.47216239e-01],
                [ -4.52068105e-02,  3.51263841e-02,  3.07536976e-02,
                4.44793489e-01, -1.63273400e-01],
                [ -2.50717045e-02,  1.70166794e-02,  4.57227166e-02,
                6.76083874e-01, -1.55256754e-01],
                [ -3.90694430e-02, -1.67460870e-02,  8.13962648e-02,
                8.36389105e-01, -9.18259738e-02],
                [ -5.14161938e-02, -3.57207133e-02,  1.10795164e-01,
                9.25123732e-01, -4.44866412e-02],
                [ -6.08589976e-02, -4.41860612e-02,  1.35794020e-01,
                9.53019912e-01, -2.21548589e-02],
                [ -4.57771169e-02, -5.25526111e-02,  1.41019853e-01,
                9.32702589e-01, -5.77563331e-07],
                [ -4.19486035e-02, -5.90387634e-02,  1.28851777e-01,
                8.87316677e-01,  1.05894354e-02],
                [ -2.47894631e-02, -7.29670549e-02,  9.41510411e-02,
                7.79023667e-01,  2.95352863e-02],
                [  2.12258453e-01,  4.65992340e-01, -6.14495953e-01,
                -2.47660051e-02,  3.66644524e-01],
                [  2.33057257e-01,  4.47057851e-01, -6.28263442e-01,
                -2.71547799e-02,  3.43419628e-01],
                [  7.85157099e-01,  4.91249250e-02,  1.44540487e-01,
                -2.05217630e-01, -1.54523367e-01],
                [  7.10324868e-01,  4.99730429e-02,  1.32239995e-01,
                -2.19171855e-01, -2.10505576e-01],
                [  8.61780950e-01,  4.35044832e-02,  1.65839097e-01,
                -1.19850814e-01,  3.16733597e-02],
                [ -2.23443268e-01,  8.46259548e-01, -4.61177284e-02,
                6.88599220e-02,  2.27742319e-01],
                [  1.43837565e-01,  9.53197441e-01,  2.27887416e-02,
                -4.57890474e-02,  1.00796460e-01],
                [  8.30286493e-01,  3.42026000e-02,  1.61106002e-01,
                -2.04570328e-01, -7.48710501e-02],

```

```
[ 7.94476579e-01, 2.83818591e-02, 1.51219549e-01,
-2.07681494e-01, -9.12497137e-02],
[ 8.11481659e-01, 4.32314894e-02, 1.43645559e-01,
-1.07778262e-01, 5.79540197e-02],
[-3.37741899e-01, 8.64927608e-01, 3.58933641e-02,
9.07183932e-02, 4.46327272e-02],
[ 5.03572673e-02, 9.35515325e-01, 1.51475395e-01,
-2.51501261e-02, -9.34471589e-02],
[ 9.78242233e-01, -3.31490295e-02, -1.05261175e-01,
4.50364242e-02, 7.37361956e-02],
[ 9.59137174e-01, -3.90023011e-02, -1.20630338e-01,
4.52591412e-02, 6.64877103e-02],
[ 8.14087183e-01, 2.23057256e-03, 7.66518490e-02,
2.02747449e-02, 1.27634826e-01],
[-4.15353978e-01, 7.18339580e-01, 3.40068066e-01,
-7.18402751e-02, -2.77950503e-01],
[ 7.64912569e-02, 7.24900638e-01, 2.74193216e-01,
-4.83952593e-02, -3.52988295e-01],
[ 9.10390848e-01, -5.36541221e-02, -4.68641899e-02,
-7.64182970e-04, 1.63870449e-01],
[ 8.73011851e-01, -5.30302303e-02, -5.89943146e-02,
-1.58989810e-03, 1.52417529e-01],
[ 7.55087669e-01, -3.56133794e-03, 5.39542523e-02,
4.24181518e-03, 2.58043480e-01],
[-1.23469883e-01, 6.07438118e-01, 6.33039205e-01,
-2.14798827e-02, 2.47973912e-01],
[-3.42866886e-01, 5.59526274e-01, 5.88212998e-01,
-2.51533511e-02, 2.18419886e-01],
[-1.60867227e-01, -1.53062640e-02, -1.57026582e-01,
1.09243755e-01, -6.61660846e-01],
[-1.37306742e-01, -2.17250612e-02, -1.58408924e-01,
1.25156181e-01, -6.71630758e-01],
[ 2.45096185e-01, -2.54584594e-02, -2.66691412e-02,
9.53148433e-02, -6.42510831e-01],
[ 2.03988660e-01, 7.85172838e-02, -3.01656215e-01,
2.28379441e-02, -6.29223343e-01],
[ 1.08926076e-01, -6.34332399e-02, -3.36565249e-02,
-9.49480406e-02, 6.81473821e-01],
[-2.63787628e-01, -6.43281164e-03, -3.58792183e-02,
-9.37962372e-02, 6.47816982e-01],
[-2.15717041e-01, -7.36588946e-02, 3.50113231e-01,
-1.95201594e-02, 6.36783774e-01],
[ 3.94306145e-01, 6.09565684e-02, 2.55337859e-01,
-2.20362094e-01, -1.84248070e-01],
[ 4.07877890e-01, 6.27256523e-02, 2.23926908e-01,
-2.10028737e-01, -1.71989219e-01],
[ 3.53156877e-01, 5.36715662e-02, 2.69603568e-01,
-2.16933220e-01, -1.80072063e-01],
[ 2.33537261e-01, -4.91732966e-02, 8.14450766e-01,
9.36688826e-02, 3.27131932e-01],
[ 2.40298209e-01, -3.38140105e-02, 8.31496963e-01,
7.52417537e-02, 2.46323612e-01],
[-6.71839466e-02, 6.58504547e-02, 5.86207684e-01,
8.72955212e-02, 9.12541435e-02],
[ 5.59835523e-02, 8.17918693e-01, -1.78458348e-01,
-1.55949404e-02, -3.34299788e-02],
[ 7.16426383e-02, 9.23428540e-01, -1.07142695e-01,
-2.78635352e-02, -4.35991157e-02],
[ 1.92496949e-01, -4.75870395e-02, 8.03173180e-01,
1.43492705e-01, 3.33862161e-01],
[ 1.87644433e-01, -3.29941019e-02, 8.58024474e-01,
1.31329947e-01, 2.55679728e-01],
[-1.02263653e-01, 6.03984268e-02, 4.72982253e-01,
7.36848375e-02, 1.12273916e-01],
```

```
[ 6.14776635e-02,  8.77962755e-01, -1.50410288e-01,
 2.20991060e-02, -4.17158213e-02],
[ 7.83728204e-02,  9.54508785e-01, -5.91095915e-02,
 1.64800945e-02, -4.32591020e-02],
[-3.24381831e-02,  1.11167162e-01,  7.84467387e-01,
-4.37718583e-02, -2.80931215e-01],
[ 1.76682383e-01,  1.90494242e-01,  5.61405509e-01,
-1.20746162e-01, -1.32570792e-01],
[-6.37386607e-02, -7.03047925e-02, -2.68934071e-01,
 1.28589792e-01,  1.88507859e-01],
[-1.56051269e-01, -7.08033944e-02, -1.45964505e-01,
 1.24253734e-01,  1.46293119e-01],
[-3.56716288e-01, -5.29910735e-02,  1.47771600e-01,
 2.87196156e-02,  1.13159584e-01],
[ 2.42173825e-01, -2.86199134e-02, -3.25958343e-02,
 1.05027814e-01, -6.55406075e-01],
[ 3.50196742e-01, -1.05016412e-02, -3.95274120e-01,
 5.92876774e-02,  2.91651778e-01],
[ 2.25671545e-01, -3.42672758e-02,  8.92876608e-01,
 1.12426805e-01,  2.67065206e-01]]])
```

Data Modeling : Linear Regression

1. Build a linear Regression model to predict the total monthly expenditure for home mortgages loan. Please refer 'deplotment_RE.xlsx'. Column `hc_mortgage_mean` is predicted variable. This is the mean monthly mortgage and owner costs of specified geographical location. Note: Exclude loans from prediction model which have NaN (Not a Number) values for `hc_mortgage_mean`.

```
In [65]: df_train.columns
```

```
Out[65]: Index(['COUNTYID', 'STATEID', 'state', 'state_ab', 'city', 'place', 'type',
'primary', 'zip_code', 'area_code', 'lat', 'lng', 'ALand', 'AWater',
'pop', 'male_pop', 'female_pop', 'rent_mean', 'rent_median',
'rent_stddev', 'rent_sample_weight', 'rent_samples', 'rent_gt_10',
'rent_gt_15', 'rent_gt_20', 'rent_gt_25', 'rent_gt_30', 'rent_gt_35',
'rent_gt_40', 'rent_gt_50', 'universe_samples', 'used_samples',
'hi_mean', 'hi_median', 'hi_stddev', 'hi_sample_weight', 'hi_samples',
'family_mean', 'family_median', 'family_stddev', 'family_sample_weight',
'family_samples', 'hc_mortgage_mean', 'hc_mortgage_median',
'hc_mortgage_stddev', 'hc_mortgage_sample_weight', 'hc_mortgage_samples',
'hc_mean', 'hc_median', 'hc_stddev', 'hc_samples', 'hc_sample_weight',
'home_equity_second_mortgage', 'second_mortgage', 'home_equity', 'debt',
'second_mortgage_cdf', 'home_equity_cdf', 'debt_cdf', 'hs_degree',
'hs_degree_male', 'hs_degree_female', 'male_age_mean',
'male_age_median', 'male_age_stddev', 'male_age_sample_weight',
'male_age_samples', 'female_age_mean', 'female_age_median',
'female_age_stddev', 'female_age_sample_weight', 'female_age_samples',
'pct_own', 'married', 'married_snp', 'separated', 'divorced',
'bad_debt', 'bins', 'pop_density', 'age_median', 'pop_bins'],
dtype='object')
```

```
In [66]: df_train['type'].unique()
type_dict={'type':{'City':1,
                  'Urban':2,
                  'Town':3,
                  'CDP':4,
                  'Village':5,
                  'Borough':6}
          }
df_train.replace(type_dict,inplace=True)
```



```
In [67]: df_train['type'].unique()
```

```
Out[67]: array([1, 2, 3, 4, 5, 6])
```

```
In [68]: df_test.replace(type_dict,inplace=True)
```

```
In [69]: df_test['type'].unique()
```

```
Out[69]: array([4, 1, 6, 3, 5, 2])
```

```
In [70]: feature_cols=['COUNTYID','STATEID','zip_code','type','pop','family_mean',
                    'second_mortgage','home_equity','debt','hs_degree',
                    'age_median','pct_own','married','separated','divorced']
```

```
In [71]: x_train=df_train[feature_cols]
         y_train=df_train['hc_mortgage_mean']
```

```
In [72]: x_test=df_test[feature_cols]
         y_test=df_test['hc_mortgage_mean']
```

```
In [73]: from sklearn.preprocessing import StandardScaler
         from sklearn.linear_model import LinearRegression
         from sklearn.metrics import r2_score, mean_absolute_error, mean_squared_error, accuracy_score
```

```
In [74]: x_train.head()
```

```
Out[74]:
```

	COUNTYID	STATEID	zip_code	type	pop	family_mean	second_mortgage	home_equity
UID								
267822	53	36	13346	1	5230	67994.14790	0.02077	0.08919
246444	141	18	46616	1	2633	50670.10337	0.02222	0.04274
245683	63	18	46122	1	6881	95262.51431	0.00000	0.09511
279653	127	72	927	2	2700	56401.68133	0.01086	0.01086
247218	161	20	66502	1	5637	54053.42396	0.05426	0.05426

```
In [75]: sc=StandardScaler()
         x_train_scaled=sc.fit_transform(x_train)
         x_test_scaled=sc.fit_transform(x_test)
```

a) Run a model at a Nation level. If the accuracy levels and R square are not satisfactory proceed to below step.

```
In [76]: linereg=LinearRegression()
         linereg.fit(x_train_scaled,y_train)
```

```
Out[76]: LinearRegression()
```

```
In [77]: y_pred=linereg.predict(x_test_scaled)
```

```
In [78]: print("Overall R2 score of linear regression model", r2_score(y_test,y_pred))
         print("Overall RMSE of linear regression model", np.sqrt(mean_squared_error(y_test,y_pred)))
```

Overall R2 score of linear regression model 0.7348210754610929
 Overall RMSE of linear regression model 323.1018894984635

The Accuracy and R2 score are good, but still will investigate the model performance at state level

b) Run another model at State level. There are 52 states in USA.

```
In [79]: state=df_train['STATEID'].unique()
state[0:5]
```

```
Out[79]: array([36, 18, 72, 20, 1])
```

```
In [80]: for i in [20,1,45]:
          print("State ID-",i)

          x_train_nation=df_train[df_train['COUNTYID']==i][feature_cols]
          y_train_nation=df_train[df_train['COUNTYID']==i]['hc_mortgage_mean']

          x_test_nation=df_test[df_test['COUNTYID']==i][feature_cols]
          y_test_nation=df_test[df_test['COUNTYID']==i]['hc_mortgage_mean']

          x_train_scaled_nation=sc.fit_transform(x_train_nation)
          x_test_scaled_nation=sc.fit_transform(x_test_nation)

          linereg.fit(x_train_scaled_nation,y_train_nation)
          y_pred_nation=linereg.predict(x_test_scaled_nation)

          print("Overall R2 score of linear regression model for state,",i,":-" ,r2_score(y_test_nation,y_pred_nation))
          print("Overall RMSE of linear regression model for state,",i,":-" ,np.sqrt(mean_squared_error(y_test_nation,y_pred_nation)))
          print("\n")
```

State ID- 20

Overall R2 score of linear regression model for state, 20 :- 0.6046603766461809
 Overall RMSE of linear regression model for state, 20 :- 307.9718899931472

State ID- 1

Overall R2 score of linear regression model for state, 1 :- 0.8104382475484617
 Overall RMSE of linear regression model for state, 1 :- 307.8275861848435

State ID- 45


Overall R2 score of linear regression model for state, 45 :- 0.788744649785525
 Overall RMSE of linear regression model for state, 45 :- 225.6961542072414

```
In [81]: residuals=y_test-y_pred
residuals
```

```
Out[81]: UID
255504    281.969088
252676    -69.935775
276314    190.761969
248614   -157.290627
286865     -9.887017
...
238088    -67.541646
242811    -41.578757
250127   -127.427569
241096   -330.820475
287763    217.760642
Name: hc_mortgage_mean, Length: 11709, dtype: float64
```

```
In [82]: plt.hist(residuals)
```

```
Out[82]: (array([6.000e+00, 3.000e+00, 2.900e+01, 7.670e+02, 7.823e+03, 2.716e+03,
        3.010e+02, 4.900e+01, 1.200e+01, 3.000e+00]),
 array([-2515.04284233, -1982.92661329, -1450.81038425, -918.69415521,
        -386.57792617, 145.53830287, 677.65453191, 1209.77076095,
        1741.88698999, 2274.00321903, 2806.11944807]),
 <BarContainer object of 10 artists>)
```




```
In [83]: sns.distplot(residuals)
```

```
/home/spx072/anaconda3/lib/python3.10/site-packages/seaborn/distributions.py:2619:
FutureWarning:
```


```
`distplot` is a deprecated function and will be removed in a future version. Please
adapt your code to use either `displot` (a figure-level function with similar fl
exibility) or `histplot` (an axes-level function for histograms).
```

```
Out[83]: <AxesSubplot:xlabel='hc_mortgage_mean', ylabel='Density'>
```



```
In [84]: plt.scatter(residuals,y_pred)
```

```
Out[84]: <matplotlib.collections.PathCollection at 0x7fc848a137f0>
```



```
In [ ]:
```